

# Modeling and Algorithms for Prof Shi's Project

Hongda Li, Yining Zhou, Xiaoping Shi

February 14, 2024

## Abstract

We propose some better algorithm for a problem in detecting structure of probability transition matrices from data.

## 1 Introduction

We describe an optimization problem introduced by Prof Shi and his student Yining. To start we define the following quantities for the optimization problem.

1.  $n \in \mathbb{N}$ . It denotes the numer of states for a Markov Chain.
2.  $p \in \mathbb{R}^{n \times n}$  denotes the probability transition matrix. It's in small case because it's also the variable for the optimization problem. It supports two types of indexing,  $p_{i,j}$  for  $i, j \in \{1, \dots, n\}$ , or  $p_j$  with  $j \in \{1, \dots, n^2\}$ . More on this later.
3.  $\eta_{ij} \geq 0$  for  $i, j \in \{1, \dots, n\}$  is a parameter of the problem. It also supports indexing with a 2D indices of type  $(i, j) \in \{1, \dots, n\} \times \{1, \dots, n\}$  as if it were a matrix. Or indexing by  $\eta_i$  for  $i \in \{1, \dots, n\}$  as if it were a matrix flattened into a vector.
4.  $\hat{p}$  is the empirically measured probability transition matrix. They are the maximal likelihood estimators for the transition probability in the transition probability matrix.
5.  $\lambda$  is the regularization parameter.
6.  $\mathbf{C}_n^m$  is the combinatoric term that counts all possible subset of size  $n, n < m$  from a super set of size  $m$ .

When  $p$  is referred to as a vector we may say  $p \in \mathbb{R}^{n^2}$ , if it's referred to as the matrix, we will use  $p \in \mathbb{R}^{n \times n}$ . When indexing  $p$  using a tuple, or a single number, it's possible to translate between the two type of indexing scheme using the following bijective map:

$$\begin{aligned}(i, j) &\mapsto k := i \times n + j \\ k &\mapsto (i, j) := (\lfloor k/n \rfloor, \text{mod}(k, n) + 1).\end{aligned}$$

We emphasize, in different programming languages and development environments, the convention of indexing a muti-array using different kind of tuples can be very different. For now we use the above indexing, which is a row major index convention (Like Python).

## 1.1 Some Mathematical Entities

We introduce some mathematical entities for more context.

## 1.2 The Optimization Problem

The optimization problem is posed as

$$\operatorname{argmin}_{p \in \mathbb{R}^{n^2}} \left\{ \sum_{i=1}^n \sum_{j=1}^n -\eta_{ij} \log(p_{i,j}) + \lambda \sum_{i=1}^{n^2} \sum_{\substack{j>i \\ p_j \neq p_i}} \frac{|p_i - p_j|}{|\hat{p}_i - \hat{p}_j|} \right\} \quad (1.2.1)$$

$$\text{s.t: } \left\{ \begin{array}{ll} \sum_{i=1}^n p_{i,j} = 1 & \forall i = 1, \dots, n \\ p_{i,j} \geq 0 & \forall i = 1, \dots, n, j = 1, \dots, n \end{array} \right\} \quad (1.2.2)$$

There are three parts to the optimization problem posed above.

1. It has a smooth differentiable function, the sum of  $\eta_{ij} \log(p_{i,j})$  but its gradient is not Lipschitz.
2. There is a non-smooth part with  $p_i - p_j$  for all indices  $1, \dots, n^2$  and  $i \neq j$ .
3. There is a polytoptic constraint set

If  $\hat{p}_i - \hat{p}_j = 0$ , it is ignored in the sum of the non-smooth part. Finally, it has a linear constraints on all  $n^2$  variables.  $p \in \mathbb{R}^{n^2}$  is a vector with the structure  $\Delta_n \oplus \Delta_n \oplus \dots \oplus \Delta_n$ . Each  $\Delta_n$  is a probability simplex. It's defined as  $\Delta_n = \{x \in \mathbb{R}_+^n : \sum_{i=1}^n x_n = 1\}$ . For simplicity we just denote using notation  $\Delta_{n \times n} = \Delta_n \oplus \Delta_n \oplus \dots \oplus \Delta_n$ .

## 1.3 Analyzing the Smooth Part of the Objective Function

The formulated problem presents challenges. Let's focus on the smooth part of the objective function. Let's denote  $f = \sum_{i=1}^{n^2} -\eta_i \log(p_i)$ . When  $\eta_{i,j} = 0$  for some integer  $1 < i < j < n$ , the domain of the objective function is not closed due to the term  $-\eta_{ij} \log(p_{i,j})$ . Here, we define that  $0 \log 0 = 0$ , as a deliberate choice.

**Claim 1.3.1** (Domain of the smooth Part is not Closed). Since the function is a sum of  $\eta_{ij} \log(p_{i,j})$ , we consider the domain of the function to be

$$\operatorname{dom} f = \bigoplus_{i=1}^{n^2} \left\{ \begin{array}{ll} \mathbb{R}_+ & \eta_i > 0 \\ \mathbb{R}_{++} & \eta_i = 0 \end{array} \right\}.$$

Observe that if there are values of  $\eta_{i,j} = 0$ , the set  $\operatorname{dom}(f)$  is not a closed set!

**Remark 1.3.1.** Adding the transition probability matrix constraint  $\Delta_{n \times n}$ , giving us  $f(x) + \delta\{\Delta_{n \times n}\}(x)$ , and the domain of the function is still not a closed set.

Similarly, the gradient of the function is the product of the derivative of different functions as well. It can be computed by

$$\begin{aligned}\nabla f(p) &= \bigoplus_{i=1}^{n^2} [-\eta_i \log(\cdot)]'(p_i) \quad \forall p \in \text{dom} f \\ &= \bigoplus_{i=1}^{n^2} \left\{ \begin{array}{ll} -\eta_i/p_i & \eta_i \neq 0 \\ 0 & \text{else} \end{array} \right\} \quad \forall p \in \text{dom} f.\end{aligned}$$

Therefore the domain of  $\nabla f = \mathbb{R}_{++}^{n^2}$  because it's the intersection of the domain  $\text{dom} f$  with the above expression we derived for the gradient of the function. It is the strictly positive quadrant of  $\mathbb{R}^{n^2}$ .

**Remark 1.3.2.** When there exists some  $\eta_i = 0$ , we have

1. The gradient  $\nabla f$  is not a globally Lipschitz Mapping.
2. And the  $\nabla^2 f$  can be a singular Hessian, and its eigenvalues are also not bounded.
3. The function is not strictly convex in the direction  $\vec{e}_i$  as well.

## 1.4 Analyzing the Probability Transition Matrix Constraints

In this section we provide a more detailed description of the transition probability matrix constraint by incorporating it with the smooth part of the objective function.

## 1.5 Existence of a Minimizer

The function may not be coersive. The smooth part has singular Hessian. The matrix inside of the one norm may be singular. Whether it is singular or not, it depends on the parameters of the problem. In the case where  $\hat{p}_i$  are all the same for  $i = 1, \dots, n^2$ , we have  $\mathbf{1}$  in the null space of matrix  $C$  because rows of matrix  $C$  are composed by row vector of form  $e_i - e_j$ . **Are there minimizers in this case?** When  $A$  has full row rank, the non-smooth term of the objective function will be coersive. Since the summed objective is a lower semi-continuous function, there exist minimizers in this case. **Is the minimier unique in this case?**

Aditinally, if all  $\eta_{i,j} \neq 0$ , then the smooth part of the objective function is a strictly convex function. It will have a unique minimizer in this case.

## 2 Modeling

The non-smooth part with the absolute value requires some amount of creativity if we were to use common optimization algorithms.

## 2.1 1-Norm Linear Compositive representation for the Non-smooth Part

**Claim 2.1.1.** The nonsmooth objective can be model as

$$\lambda \sum_{i=1}^{n^2} \sum_{j>i}^{n^2} \frac{|p_i - p_j|}{|\hat{p}_i - \hat{p}_j|} = \|Cp\|_1 \quad \text{where } C \text{ is } \mathbf{C}_2^{n^2} \text{ by } n^2.$$

The one norm represents the summation of absolute values. The transformation  $Cp$  is linear transformation and it's a vector of length  $\mathbf{C}_2^{n^2}$ . The vector is long and it has a dimension of  $(1/2)(1 + n^2)n^2$ . Each term inside of the summation is a row of the matrix  $C$ . Each row of matrix  $C$  has exactly 2 non-zero elements in it. Suppose that  $i \in \{1, \dots, \mathbf{C}_2^{n^2}\}$  denoting the index for a specific row of matrix  $C$  denotes the index of a specific row, and  $j \in \{1, \dots, n^2\}$  denotes a specific column of matrix  $C$ . Mathematically describing the matrix is difficult, but it can be algorithmically defined. A sparse matrix format can be described by as a mapping from  $(i, j)$ , the set of all indices to the element in the vector. The following [algorithm 1](#) construct such a mapping.

---

### Algorithm 1 Matrix Make Algorithm

---

```

1: Let  $C$  be a  $\mathbf{C}_2^{n^2}$  by  $n^2$  zero matrix.
2: for  $i = 1, \dots, n^2$  do
3:   for  $j = 1, \dots, n^2$  do
4:     if  $|\hat{p}_i - \hat{p}_j| == 0$  then
5:       continue
6:     end if
7:      $C[i \times n^2 + j, i] := \lambda / |\hat{p}_i - \hat{p}_j|$ 
8:      $C[i \times n^2 + j, j] := -\lambda / |\hat{p}_i - \hat{p}_j|$ 
9:   end for
10: end for
```

---

**Remark 2.1.1.** In practice, we should use sparse matrix data format such as SCR, SCC in programming languages.

## 2.2 Modeling it for Sequential Quadratic Programming

To start, refer to [Sequential Qaudratic Programming Wikipedia](#) for brief overview about what sequential quadratic programming problem is. To use sequential programming, we model the non-smooth  $\|Cp\|_1$  parts of the objective function as a linear constraints. Define  $u \in \mathbb{R}_+^{\mathbf{C}_2^{n^2}}$

then the below problem is equivalent to the original formulation:

$$\underset{p,u}{\operatorname{argmin}} \left\{ \sum_{i=1}^n \sum_{j=1}^n -\eta_{ij} \log(p_{i,j}) + \sum_{i=1}^{\mathbf{C}_2^{n^2}} u_k \right\} \quad (2.2.1)$$

$$\text{s.t.: } \left\{ \begin{array}{l} -u \leq Cp \leq u \\ p \in \Delta_{n \times n} \\ u \in \mathbb{R}_+^{\mathbf{C}_2^{n^2}} \end{array} \right\} \quad (2.2.2)$$

This is a Non-linear programming problem and it has a convex objective. Common NLP packages in programming languages can solve this efficiently. However it's potentially possible that these solvers are not adapted to huge sparse matrix  $C$  that has special structure to it.

**Remark 2.2.1.** To formulate into a linear programming with some relaxations, consider that the non-linear objective  $-\eta_{ij} \log(p_{i,j})$  can be discretized.

## 2.3 A Quick Tour into Operator Splitting Method

Operator splitting method is relatively new and doesn't have a mature toolings in conventional software yet. Some julia implementation for splitting method exists. These type of method requires the use of the Proximal Mapping, which can be non-trivial to implement for even simple functions. We refer readers to [1, section 2.7] of Ryu's Textbook for more information on the Douglas Rachford method.

## 2.4 Modeling it for Douglas Rachford

The Douglas Rachford algorithms solves objective function of the type  $f + g$  where  $f, g$  are both convex function and they are proximal friendly. It has the assumption  $\text{ri.dom}(f) \cap \text{ri.dom}(g) \neq \emptyset$ . Recall that proximal operator of function  $f$  is the resolvent operator on the subgradient  $(I + \partial f)^{-1}$ . Subgradient is a generalize type of gradient that can handle continuous function that is not necessarily differentiable. To model the original formulation in such a form, we introduce these quantities:

1.  $f_1(x_1) : \mathbb{R}^{n^2} \mapsto \mathbb{R} := \sum_{i=1}^{n^2} -\eta_{i,j}(\log(p_{i,j}))$
2.  $f_2(x_2) : \mathbb{R}^{\mathbf{C}_2^{n^2}} \mapsto \mathbb{R} := \|x_2\|_1.$
3.  $f_3(x_3) : \mathbb{R}^{n^2} \mapsto \bar{\mathbb{R}} := \delta_{\Delta_{n \times n}}(x_3)$

The above three functions represent the three parts of the summed objective.

Let  $f(x_1, x_2, x_3) := f_1(x_1) + f_2(x_2) + f_3(x_3)$ . However, they share different variables  $x_1, x_2, x_3$ , from different dimension. We want  $g$  to represents the constraints that  $x_2 = Cp$ , and  $x_1 = x_3 = p$ . Simplifying:  $x_2 = Cx_1, x_1 = x_3$ . This is a linear system of the form

$$\begin{bmatrix} C & -I & \mathbf{0} \\ I & \mathbf{0} & -I \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \mathbf{0}.$$

Simply denote the above as  $Ax = \mathbf{0}$ , then  $g$  is just a non-smooth function that happens to be an indicator function of a convex set. The convex set is the set of all solutions to the linear system. Therefore setting  $g = \delta_{Ax=\mathbf{0}}(x)$ , to be the indicator of the linear constraints, we had a complete equivalent representation of the original form of the problem.

The matrix  $A$  is a  $(\mathbf{C}_2^{n^2} + n^2) \times (3n^2)$  matrix. It's still a sparse matrix. The matrix would have to be encoded carefully by translating copies of the same matrix to 3 different locations.

## Modeling it for Chambolle Pock

The method of Chmabolle Pock has better convergence theory. It also comes with an acceleration scheme when certain conditions for the objective function applies. We refer readers to Chambolle and Pock's paper [2] for more information about the algorithm. For the modeling purpose in this section, it's sufficient to understand the the Chambolle Pock algorithm solves the set-valued equation  $\mathbf{0} \in \partial f(x) + \partial g(Lx)$  where  $f, g$  are convex proper and closed function and  $\text{ri.dom} f \cap \text{ri.dom} g \neq \emptyset$ . And  $L$  is a bounded linear operators. In our case, the solution of the above equation yield the minimizers of the objective  $f(x) + g(Lx)$ . We claim that our problem can be put into the same format that is accepted by the Chambolle Pock algorithm. Consider defining the following quantities

1. Let  $\phi(z)$  be a function mapping from  $\mathbb{R}^{n^2} \oplus \mathbb{R}^{\mathbf{C}_2^{n^2}}$  such that the vector  $z$  admits representation  $[x \ y]^T$  where  $x \in \mathbb{R}^{n^2}$  and  $y \in \mathbb{R}^{\mathbf{C}_2^{n^2}}$ . This function handles the smooth and nonsmooth objective given as

$$\phi(z) = f(x) + \|y\|_1.$$

2. Define

$$L = \begin{bmatrix} I & \mathbf{0} \\ A & -I \end{bmatrix}$$

3. Let  $\Phi(z) = \delta\{\Delta_{n \times n}\}(x) + \delta\{\mathbf{0}\}(y)$ .

## 2.5 Another Method that is not Cursed by the Dimension

## 3 Implementation, Software, and Toolings

The software for Sequential Quadratic programming is relatively mature. Whether they support sparse matrix inversion depends on the developers. The software operator splitting method is less common. To use Operator splitting method such as Douglas Rachford (Equivalently the ADMM via a the Dual), Peachman Rachford, it's required to compute the proximal operator of functions. The following theorem is relevant to efficiently compute the proximal mapping for our model.

**Theorem 1** (Separable Proximal Theorem). Suppose that function  $f(x)$  can be written as the Euclidean cross product  $(f_1(x_1), f_2(x_2), \dots, f_m(x_m))$ , and we assume that each of

$x_1, x_2, \dots, x_n$  are vector in  $\mathbb{R}^{k_1}, \mathbb{R}^{k_2}, \dots, \mathbb{R}^{k_m}$  such that  $x = \bigoplus_{i=1}^m x_i$ , then we can write the proximal gradient operator of  $f$  in parallel form which is given as:

$$\text{prox}_{f,\lambda}((x_1, x_2, \dots, x_m)) = (\text{prox}_{f_1,\lambda}(x_1), \text{prox}_{f_2,\lambda}(x_2), \dots, \text{prox}_{f_m,\lambda}(x_m)).$$

In general, if there exists any permutation of the list of  $x_1, x_2, \dots, x_m$ , the parallelization property of prox will be come applicable.

Recall objective  $f(x_1, x_2, x_3) = f_1(x_1) + f_2(x_2) + f_3(x_3)$  from section 2.1. It has the applicable form and hence we have:

$$\text{prox}_{\lambda f}((x_1, x_2, x_3)) = \text{prox}_{\lambda f_1}(x_1) \oplus \text{prox}_{\lambda f_2}(x_2) \oplus \text{prox}_{\lambda f_3}(x_3).$$

Next, we proceed to write each  $f_i$  for  $i = 1, 2, 3$  into separable sum as well.  $f_1, f_2$  are trivial to do however  $\delta_{\Delta_{n \times n}}$  requires some thought. Consider  $x_3 = \bigoplus_{i=1}^n x_{3,i}$ , so that the vector is the direct product, then

$$\delta_{\Delta_{n \times n}}(x_3) = \sum_{i=1}^n \delta_{\Delta_n}(x_{3,i}).$$

The indicator function  $\delta_{\Delta_{n \times n}}$  acts on groups of  $n$  elements in the vector  $x_3$  individually and separately, hence it admits the above separable summation form. We now list the proximal mapping for each of the atomic operations:  $\lambda \log(\cdot)$ ,  $\lambda |\cdot|$ ,  $\delta_{\Delta_n}$  and  $\delta_{Ax=0}(\cdot)$ . Most of them can be found in the literatures.

$$\begin{aligned} \text{prox}_{-\lambda \log(\cdot)}(x) &= \frac{x + \sqrt{x^2 + 4\lambda}}{2} \\ \text{prox}_{\lambda |\cdot|}(x) &= \text{sign}(x) \max(|x| - \lambda, 0) \\ \text{prox}_{\delta_{\Delta_n}} &= \Pi_{\Delta_n} \\ \text{prox}_{\delta_{\{Ax=0\}}} &= \Pi_{Ax=0}. \end{aligned}$$

### 3.1 Projection onto Probability Simplex and Sparse Linear System

Projecting onto a probability simplex and a linear system that is sparse requires a bit more attention for software implementation sake. We introduce the following theorem from Beck textbook to assist with the projection onto a probability simplex.

**Theorem 2** (Projecting onto Hyperplane Box Intersections). Let  $C \subseteq \mathbb{R}^n$  to be non-empty and defined as  $C := \text{Box}[l, u] \cap H_{a,b}^-$ , where  $\text{Box}[l, u] = \{x \in \mathbb{R}^n : l \leq x \leq u\}$ . We assume that  $l \in [-\infty, \infty)$  and  $u \in (-\infty, \infty]$ . And  $H_{a,b}^- = \{x \in \mathbb{R}^n : \langle a, x \rangle = b\}$ , where  $a \in \mathbb{R}^n \setminus \{0\}$ . Then the projection onto the set is equivalent to the following conditions

$$\begin{aligned} \Pi_C(x) &= \Pi_{\text{box}[l,u]}(x - \mu a) \\ \text{where } u &\in \mathbb{R} \text{ solves: } \langle a, \Pi_{\text{Box}[l,u]}(x - ua) \rangle = b. \end{aligned}$$

*Proof.* See [3, theorem 6.27] by Beck for more information. □

**Remark 3.1.1.** The equation for  $u$  has a continuous mapping  $\mathbb{R} \mapsto \mathbb{R}$ , it can be solved efficiently using method such as the bisection method.

**Theorem 3** (Projecting onto the Probability Simplex). The projection onto the simplex  $\Delta_n$  can be computed via

$$\begin{aligned}\Pi_{\Delta_n}(x) &= \Pi_{\mathbb{R}_+}(x - \mu a) \\ \text{where } \mu \in \mathbb{R} \text{ solves: } \langle \mathbf{1}, \Pi_{\mathbb{R}_+}(x - \mu \mathbf{1}) \rangle &= 1.\end{aligned}$$

*Proof.* Use the previous theorem 2, observe that  $u$  can be a vector of infinity and hence setting  $\mathbb{R}_+$  to be the box, and the probability simplex constraints  $\sum_{i=1}^n x_i = 1$  is the hyperplane  $H_{a,b}^-$  with  $a = \mathbf{1}, b = 1$ .  $\square$

**Theorem 4** (Projection onto Affine Spaces). Define the set  $C = \{x \in \mathbb{R}^n : Ax = b\}$ , assuming that  $A \in \mathbb{R}^{m \times n}$  is full rank, then the projection onto the set is computed via

$$\Pi_C(x) = x - (A^T A)^{-1} A(Ax - b).$$

**Remark 3.1.2.** In the case when  $A$  is not full rank, do a reduce QR decomposition of  $A = QR$  where  $R$  is  $k \times n$  with  $k < m$ . In the case of our application, the matrix  $A$  is a full rank matrix. The matrix  $(A^T A)^{-1}$  in our case can be very slow to invert since it's  $\mathbf{C}_2^{n^2} \times \mathbf{C}_2^{n^2}$ . One way to make it easier is to use sparse matrix invert method. In our case the matrix we are inverting is sparse and symmetric, it can be solved efficiently say for example,

## 4 The Software APIs for the Programmers

Option [SLSQP](#) in “`scipy.optimize.minimize`” performs sequential quadratic programming. I am not sure whether they would support efficient sparse matrix inverse. [Interior points method](#) (the engine behind many constraint optimization solver) is also a good way of solving the problem, it's also more specialized than the SLSQP solver. However their availability in python seems to be limited. See [here](#) for tutorials about how to use “`scipy.optimize`” module. Visits [pyprox](#) for a package that implements various type of proximal operators in python.

### 4.1 Scipy.optimize.SLSQP

[scipy optimization documentations with examples](#). The “`scipy.SLSQP`” programming API accepts programming problem of the following form

$$\begin{aligned}\min_{l \leq x \leq u} f(x) \text{ s.t:} \\ \begin{cases} c_j(x) = 0, & j \in \Upsilon \\ g_j(x) \geq 0, & j \in \mathcal{I} \end{cases}\end{aligned}$$

For each element  $i = n, \dots, N$ , we have the constraints in the form of  $l_i \in [-\infty, \infty)$  and  $u_i \in (-\infty, \infty]$ . The programmer is required to provide

1.  $\nabla f(x)$ , the gradient of the objective and the objective function in the form of a function handle in python.



2. The gradient and function of the equality and the inequality constraint of the problem.

Representing the sequence of constraints  $c_j(x)$  for  $J \in \Upsilon$  as a vector function  $C(x)$ . Similarly, we represent inequality constraints as the vector function  $G(x)$ . Recall formulation from [section 2.2](#). The inequality constraints are  $-u \leq Cp \leq u$ ,  $p \in \Delta_{n \times n}$ . The box constraint is  $u \in \mathbb{R}_+^{\mathbf{C}_2^{n^2}}$ . The equality and inequality constraints can be modeled by the linear system:

$$\underbrace{\begin{bmatrix} -C & I \\ C & I \end{bmatrix}}_{=: \mathcal{C}} \begin{bmatrix} p \\ u \end{bmatrix} \geq \mathbf{0}$$

$$\underbrace{\begin{bmatrix} \mathbf{1}_n^T & & & \\ & \mathbf{1}_n^T & & \\ & & \ddots & \\ & & & \mathbf{1}_n^T \end{bmatrix}}_{\mathcal{G}} p - \mathbf{1} = \mathbf{0}.$$

Here  $\mathbf{1}_n \in \mathbb{R}^n$  is a vector of 1. Observe that,  $\mathcal{C}$  is a  $2\mathbf{C}_2^{n^2}$  by  $2n + \mathbf{C}_2^{n^2}$  matrix.  $\mathcal{G}$  is a  $n \times n^2$  matrix. The matrix  $\mathcal{C}, \mathcal{G}$  denoted above are the Jacobi of the inequality and equality constraints respectively. The gradient of the objective is stated back in [section 1.3](#). At this point, we have everything ready for coding it up in python.

## References

- [1] E. K. Ryu and W. Yin, *Large-Scale Convex Optimization: Algorithms & Analyses via Monotone Operators*. Cambridge: Cambridge University Press, 2022. [Online]. Available: <https://large-scale-book.mathopt.com/>
- [2] A. Chambolle and T. Pock, “A first-order primal-dual algorithm for convex problems with applications to imaging,” Jun. 2010. [Online]. Available: <https://hal.science/hal-00490826>
- [3] A. Beck, *First-Order Methods in Optimization* / *SIAM Publications Library*, ser. MOS-SIAM Series in Optimization. SIAM. [Online]. Available: <https://epubs.siam.org/doi/book/10.1137/1.9781611974997>